

Summary

“BallPath” project uses A* search algorithm for pathfinding. This algorithm make it possible to search optimal (shortest) path. A* is *complete* and will always find a solution if one exists.

Algorithm description

A* uses a [best-first search](#) and finds a least-cost path from a given initial [node](#) to one [goal node](#) (out of one or more possible goals).

It uses a distance-plus-cost [heuristic](#) function (usually denoted $f(x)$) to determine the order in which the search visits nodes in the tree. The distance-plus-cost heuristic is a sum of two functions:

- the path-cost function, which is the cost from the starting node to the current node (usually denoted $g(x)$)
- and an [admissible](#) "heuristic estimate" of the distance to the goal (usually denoted $h(x)$).

The $h(x)$ part of the $f(x)$ function must be an [admissible heuristic](#); that is, it must not overestimate the distance to the goal. Thus, for an application like [routing](#), $h(x)$ might represent the straight-line distance to the goal, since that is physically the smallest possible distance between any two points or nodes.

If the [heuristic](#) h satisfies the additional condition $h(x) \leq d(x, y) + h(y)$ for every edge x, y of the graph (where d denotes the length of that edge), then h is called [monotone, or consistent](#). In such a case, A* can be implemented more efficiently—roughly speaking, no node needs to be processed more than once (see *closed set* below)—and A* is equivalent to running [Dijkstra's algorithm](#) with the [reduced cost](#) $d'(x, y) = d(x, y) - h(x) + h(y)$.

Algorithm complexity (speed of search and worst case)

The [time complexity](#) of A* depends on the heuristic. In the worst case, the number of nodes expanded is [exponential](#) in the length of the solution (the shortest path), but it is [polynomial](#) when the search space is a tree, there is a single goal state, and the heuristic function h meets the following condition:

$$|h(x) - h^*(x)| = O(\log h^*(x)) ,$$

where h^* is the optimal heuristic, the exact cost to get from x to the goal. In other words, the error of h will not grow faster than the [logarithm](#) of the “perfect heuristic” h^* that returns the true distance from x to the goal.

Implementation notes

As heuristic cost estimate function used [Manhattan length](#): $l_M = |x_1 - x_2| + |y_1 - y_2|$.

To prevent a premature pessimisation BallPath implementation of A* uses next data structures (see `src/core/pathfinder.h`):

- for *open list*: `std::priority_queue`.
Rationale: fast retrieving of minimal element; convenient usage for open list.
- for *closed list*: `std::unordered_set` (C++11).
Rationale: fast element pushing and finding; convenient usage for closed list.